

CV201, HW #3

Oren Freifeld and Meitar Ronen
The Department of Computer Science, Ben-Gurion University

Release Date: 7/12/2019
Submission Deadline: 9/1/2020, 20:00

Abstract

Please make sure you have carefully read the general instructions in the course website at https://www.cs.bgu.ac.il/~cv201/HW_Instructions. These instructions address, among other things, what to do or not do with figures, Jupyter Notebooks, *etc.*, as well as what compression files are legit or not.

As for this assignment: We will test the Gibbs sampler by computing expectations and comparing with the exact results (which we obtained using dynamic programming) from HW#2. We will also pretend that the Ising model is a good model for images, and study image restoration from Ising-model samples corrupted by noise.

Version Log

- 1.02, 26/12/2019. in Computer Exercise 2: The maximum-likelihood estimate is of course $\arg \max_x p(y|x)$ (and not $\arg \max_x p(x|y)$ as it was written there before).
- 1.01, 7/12/2019. Fixed deadline (from Jan 2019 to Jan 2020)
- 1.00, 7/12/2019. Initial release.

Contents

1	Reminder: Gibbs Sampling from the Ising-Model Prior	2
2	Reminder: Gibbs Sampling from the Ising-Model Posterior, Assuming Gaussian IID Additive Noise	2
3	Estimation: Comparisons with the Nearly-exact Values	3
4	Image Restoration	5

Remark 1 *Before you start coding a single line in this assignment, do yourself a favor and read the entire document.* \diamond

1 Reminder: Gibbs Sampling from the Ising-Model Prior

Recall the Ising-model prior, $p(x)$, is

$$p(x) \propto \exp\left(\frac{1}{\text{Temp}} \sum_{s \sim t} x_s x_t\right) \quad \text{Temp} > 0, \quad (1)$$

and that

$$\begin{aligned} p(x_s | s x) &= \frac{p(x_s, s x)}{p(s x)} = \frac{p(x)}{p(s x)} = \frac{\frac{1}{Z_{\text{Temp}}} \exp\left(\frac{1}{\text{Temp}} \sum_{s \sim t} x_s x_t\right)}{\sum_{x'_s} \frac{1}{Z_{\text{Temp}}} \exp\left(\frac{1}{\text{Temp}} \sum_{s \sim t} x'_s x_t\right)} \\ &= \frac{\exp\left(\frac{1}{\text{Temp}} \sum_{t: t \in \eta_s} x_s x_t\right)}{\sum_{x'_s} \exp\left(\frac{1}{\text{Temp}} \sum_{t: t \in \eta_s} x'_s x_t\right)} = \frac{\exp\left(\frac{1}{\text{Temp}} \sum_{t: t \in \eta_s} x_s x_t\right)}{\exp\left(-\frac{1}{\text{Temp}} \sum_{t: t \in \eta_s} x_t\right) + \exp\left(\frac{1}{\text{Temp}} \sum_{t: t \in \eta_s} x_t\right)}. \end{aligned} \quad (2)$$

Thus,

$$p(x_s | s x) = p(x_s | x_{\eta_s}) \quad (3)$$

and

$$\begin{cases} p(x_s = 1 | s x) = \frac{\exp\left(\frac{1}{\text{Temp}} \sum_{t: t \in \eta_s} x_t\right)}{\exp\left(-\frac{1}{\text{Temp}} \sum_{t: t \in \eta_s} x_t\right) + \exp\left(\frac{1}{\text{Temp}} \sum_{t: t \in \eta_s} x_t\right)} \propto \exp\left(\frac{1}{\text{Temp}} \sum_{t: t \in \eta_s} x_t\right) \\ p(x_s = -1 | s x) = \frac{\exp\left(-\frac{1}{\text{Temp}} \sum_{t: t \in \eta_s} x_t\right)}{\exp\left(-\frac{1}{\text{Temp}} \sum_{t: t \in \eta_s} x_t\right) + \exp\left(\frac{1}{\text{Temp}} \sum_{t: t \in \eta_s} x_t\right)} \propto \exp\left(-\frac{1}{\text{Temp}} \sum_{t: t \in \eta_s} x_t\right) \end{cases} \quad (4)$$

So sampling from this conditional distribution is just flipping a coin, where the bias on the coin is affected by the states of the neighbors of x_s .

2 Reminder: Gibbs Sampling from the Ising-Model Posterior, Assuming Gaussian IID Additive Noise

Under this assumption, letting x denote the sample from the Ising-model prior and y denote the noise, we have:

$$p(x|y) \propto \exp\left(\frac{1}{\text{Temp}} \left(\sum_{s \sim t} x_s x_t\right) - \frac{1}{2\sigma^2} \sum_s (y_s - x_s)^2\right) \quad (5)$$

and

$$\begin{aligned} p(x_s | s x, y) &= \frac{p(x_s, s x | y)}{p(s x | y)} = \frac{p(x | y)}{p(s x | y)} \\ &= \frac{\exp\left(\frac{1}{\text{Temp}} \left(\sum_{s \sim t} x_s x_t\right) - \frac{1}{2\sigma^2} \sum_s (y_s - x_s)^2\right)}{\sum_{x'_s} \exp\left(\frac{1}{\text{Temp}} \left(\sum_{s \sim t} x'_s x_t\right) - \frac{1}{2\sigma^2} \sum_s (y_s - x'_s)^2\right)} \\ &\propto \exp\left(\frac{1}{\text{Temp}} \left(\sum_{t: t \in \eta_s} x_s x_t\right) - \frac{1}{2\sigma^2} (y_s - x_s)^2\right). \end{aligned} \quad (6)$$

Thus,

$$p(x_s | x, y) = p(x_s | x_{\eta_s}, y) = p(x_s | x_{\eta_s}, y_s) \quad (7)$$

and

$$p(x_s = 1 | x, y) \propto \exp \left(\frac{1}{\text{Temp}} \left(\sum_{t:t \in \eta_s} x_t \right) - \frac{1}{2\sigma^2} (y_s - 1)^2 \right) \quad (8)$$

$$p(x_s = -1 | x, y) \propto \exp \left(\frac{1}{\text{Temp}} \left(- \sum_{t:t \in \eta_s} x_t \right) - \frac{1}{2\sigma^2} (y_s + 1)^2 \right). \quad (9)$$

Again, the sampling is nothing more than flipping a coin, where here the bias on the coin is affected by both (i) the states of the neighbors of x_s and (ii) y_s , the measurement associated with x_s .

3 Estimation: Comparisons with the Nearly-exact Values

In HW #2, we used exact sampling to estimate

$$E_{\text{Temp}}(X_{(1,1)}X_{(2,2)}) \text{ and } E_{\text{Temp}}(X_{(1,1)}X_{(8,8)})$$

through their empirical expectations,

$$\hat{E}_{\text{Temp}}(X_{(1,1)}X_{(2,2)}) \triangleq \frac{1}{10000} \sum_{n=1}^{10000} x_{(1,1)}(n)x_{(2,2)}(n) \quad (10)$$

$$\hat{E}_{\text{Temp}}(X_{(1,1)}X_{(8,8)}) \triangleq \frac{1}{10000} \sum_{n=1}^{10000} x_{(1,1)}(n)x_{(8,8)}(n) \quad (11)$$

where $x(1), x(2), \dots, x(10000)$ were 10,000 exact samples from the Ising model. We refer to these expectations as “nearly-exact” for the following reason: the term “exact” stems from the fact the samples were obtained using exact sampling, while the “nearly” qualifier is added since an empirical expectation, *i.e.*, the average over the samples (*i.e.*, AKA the sample mean) is a random quantity which, by the Law of Large Numbers, converges to the true expectation as the number of samples tends to infinity.

Now, in the current HW assignment, instead of using exact sampling, we will produce the samples using MCMC, particularly the Gibbs-sampling method.

Computer Exercise 1 *Build a Gibbs sampler for the 8×8 Ising model at temperatures $\text{Temp} \in \{1, 1.5, 2\}$. Compute the empirical expectations of*

$$E_{\text{Temp}}(X_{(1,1)}X_{(2,2)}) \text{ and } E_{\text{Temp}}(X_{(1,1)}X_{(8,8)}) \quad (12)$$

using two different methods; see below. Compare the estimates to the nearly-exact values computed in the previous HW assignment; see the table below.

Programming Note: *A convenient way to handle boundaries is to embed the $n \times n$ lattice in the interior of an $(n+2) \times (n+2)$ lattice whose boundary values are set to zero, and then visit only the $n \times n$ interior sites of the larger lattice. \diamond*

Method 1. **Independent Samples.** At each temperature, draw 10,000 (approximated) samples from the Ising-model prior,

$$p(x) \propto \exp \left(\frac{1}{\text{Temp}} \sum_{s \sim t} x_s x_t \right), \quad (13)$$

using the Gibbs sampler. For each sample, initiate the Gibbs sampler at a random configuration; *i.e.*, sample all the pixels *i.i.d.* using a fair coin, and assign the values $\{-1, 1\}$ accordingly. One way to do it in Python is using

`np.random.randint(low=0,high=2,size=(8,8))*2-1`.

Next, update sites in a deterministic, raster-scan, order (a fixed site-visitation schedule). The sample is the obtained configuration after 25 passes through the entire graph (*i.e.*, after 25 “sweeps” – where each sweep involves $8 \times 8 = 64$ single-site updates). Putting it differently, you are asked to run 10,000 such Markov Chains, where each chain has 25×64 iterations. To be clear, the initialization should always be done using *i.i.d.* coin flips as mentioned above, but please use a different realization of that random initialization for each Markov chain. Let the obtained 10,000 (approximated) samples be denoted by $x(1), \dots, x(10000)$. Your estimates are then

$$\frac{1}{10000} \sum_{n=1}^{10000} x_{(1,1)}(n)x_{(2,2)}(n) \text{ and } \frac{1}{10000} \sum_{n=1}^{10000} x_{(1,1)}(n)x_{(8,8)}(n). \quad (14)$$

Method 2. **Ergodicity.** Beginning with a random configuration (*i.e.*, sampling all the pixels *i.i.d.* using a fair coin, as before), run the Gibbs sampler for 25,000 sweeps of the lattice (*i.e.*, a single Markov Chain of $25,000 \times 64$ iterations). Use the empirical averages of $x_{(1,1)}x_{(2,2)}$ and $x_{(1,1)}x_{(8,8)}$ over all but the first 100 sweeps (*i.e.*, following a so-called “burn-in period” of 100 sweeps, use the 24,900 configurations obtained at the end of each of the remaining sweeps to compute the empirical averages).

Remark 2 *Note this method is related to third part of the theorem on MCMC from the lecture slides.* \diamond

Remark 3 *Note that for computing the average of, say, such 24,900 quantities, you do not need to store 24,900 values in memory. In effect, suppose you have a sequence of N values, denoted by z_1, z_2, \dots, z_N and you want to compute $\frac{1}{N} \sum_{n=1}^N z_n$. Then you can discard each z_i once you iteratively update the current*

temperature	$\hat{E}_{\text{Temp}}(X_{(1,1)}X_{(2,2)})$	$\hat{E}_{\text{Temp}}(X_{(1,1)}X_{(8,8)})$
1	0.95	0.9
1.5	your result from HW2	your result from HW2
2	your result from HW2	your result from HW2

Table 1: Expectations estimated from exact samples using dynamic programming

estimate of the mean:

$$\mu^{[1]} = z_1 \quad (15)$$

$$\mu^{[n]} = \frac{1}{n} \sum_{m=1}^n z_m = \frac{1}{n} \left(\left(\sum_{m=1}^{n-1} z_m \right) + z_n \right) \quad (16)$$

$$= \frac{1}{n} \left(\frac{n-1}{n-1} \left(\sum_{m=1}^{n-1} z_m \right) + z_n \right) = \frac{(n-1)\mu^{[n-1]} + z_n}{n} \quad n \in 2, 3, \dots, N. \quad (17)$$

◇

Exercise 1 Compare the results obtained from the two methods, and also compare these results with the values from exact sampling (obtained using dynamic programming). Speculate about the reasons for any substantial discrepancies. ◇

4 Image Restoration

These experiments are to be performed on a 100×100 lattice.

Computer Exercise 2 At each of the three temperatures $\text{Temp} \in \{1, 1.5, 2\}$:

1. Using Gibbs sampling, generate a single 100×100 sample, x , from the 100×100 Ising-model prior using 50 raster-scan sweeps.
2. Add to the sample an array of 100×100 Gaussian noise values, sampled IID from $\mathcal{N}(0, 2^2 = 4)$. We denote this noise array by η . Here is one way to achieve this in Python:

```
eta = 2*np.random.standard_normal(size=(100,100))
y = x + eta
```

3. Using Gibbs sampling, generate a sample from the posterior distribution of the uncorrupted sample given the corrupted sample,

$$x \sim p(x|y), \quad (18)$$

again using 50 sweeps. To remove any doubt, even though the y values are continuous, the x values are still discrete. Thus, even when you are using the Gibbs sampler for sampling from $p(x|y)$, you should still initialize your x values using random i.i.d. flips of a fair coin, to produce values in $\{-1, 1\}$. Particularly, do **not** initialize x by setting it to y .

4. For comparison, compute the “ICM” (Iterated Conditional Modes) restoration: upon a random initialization, visit sites, in a raster order, and replace at each site s the current value of x_s by its most-likely value (of the two possibilities, ± 1), under the posterior distribution and conditioned on the four neighbors:

$$x_s^{new} = \arg \max_{x_s \in \{-1, 1\}} p(x_s | x, y) \quad (19)$$

where

$$p(x_s | x, y) = p(x_s | x_{\eta_s}, y) = p(x_s | x_{\eta_s}, y_s). \quad (20)$$

This “greedy” algorithm will converge within a few raster cycles.

5. For yet another comparison, display the maximum-likelihood estimate, $\arg \max_x p(y|x)$. Recall that each pixel in x takes values in $\{-1, 1\}$; i.e., the maximum-likelihood estimate should also take only these discrete values. Particularly, it is **not** simply (the continuous) y .
6. Display all five images (from the previous 5 steps) on a single plot. (see the subplot and imshow commands). When showing x or its estimates, set the imshow optional parameters vmin and vmax to -1 and $+1$, respectively. When showing y , however, do not do that since it may contain values outside the range $[-1, 1]$. \diamond

Remark 4 Note well: the ICM result is (usually) **not** the argmax of the posterior. **So usually ICM is not the MAP solution** (where MAP stands for the maximum-a-posteriori); rather, ICM converges to a usually-local maximum of the posterior. Also, just like in the case of Gibbs sampling, when one can apply it to distributions that are not Gibbs distributions (i.e., not MRFs) but using Gibbs sampling to Gibbs distributions is especially convenient (due to the structure of the conditionals), we have a similar situation for ICM: one can also apply ICM to distributions that are not Gibbs distributions, but using it for Gibbs distributions will be especially convenient, and for the same reason. \diamond